

Closing the Gap on Personal Writing Style: An Agentic Six-Day Companion to Baumler et al., and Why AI-Text Detection Is an Arms Race

Anonymous Author(s)¹

Withheld for blind review

Abstract. Baumler et al. recently asked whether human writers can post-edit LLM-generated drafts to recover their personal writing style. Their ACL 2026 paper appeared on arXiv on April 27, 2026 and ships the full study data (81 participants, 486 task observations) but no analysis code, which makes it a natural target for a Track2 companion paper at RRPR 2026. Using an agentic-research harness in the spirit of Zaiss et al., we reproduce, critique, and extend the design *in six days end-to-end* (April 27 – May 3, 2026), with the human investigator acting only as a reviewer-in-the-loop. We re-implement the full LUAR-MUD / permutation-test / rmcrr pipeline from the released logs, reproduce all seven preregistered hypotheses, and reproduce the headline repeated-measures correlation between perceived self-similarity and embedding-measured self-similarity to three decimal places ($r = +0.244$, $p < 10^{-8}$, $n = 648$). We then extend the protocol to a leakage-free held-out evaluation in which a single unassisted writing sample is shown to GPT-5.5 and ClaudeOpus 4.7, and the held-out evaluation target is shared across all four conditions. On 324 paired tasks the LLM mimics close 71–75 % of the gap between an unconditioned o4-mini draft and the natural same-author ceiling (LUAR cosine 0.701); the human post-edit closes 24 %. A Friedman omnibus ($\chi^2=419.3$, $p=1.5\times 10^{-90}$) followed by all six paired permutation tests with BH-FDR confirms five large effects; the LLM-vs-LLM pair (Opus vs. GPT-5.5) is statistically tied; both LLMs beat the human post-edit on ~ 80 % of tasks. Finally, we cast the same data as an arms race: a leave-authors-out linear-SVM detector on LUAR embeddings reaches AUC 0.999 on the unconditioned o4-mini draft, 0.967 on the human post-edit, 0.942 on Opus 4.7, and 0.914 on GPT-5.5. Detectability drops monotonically as the approach gets closer to the participant’s natural style, but no approach reaches chance: a residual stylometric signature persists even in the strongest mimic. All code, all 648 mimic drafts, the trained detectors, and the agentic harness are released.

Keywords: Reproducible research · Authorship style embeddings · LLM evaluation · AI-text detection · Permutation tests · Held-out protocols · Agentic research.

1 Introduction

Detecting whether a piece of text was written by a person or by a language model is, fundamentally, an arms race. Every new generation of large language model raises the bar for the detector by closing some gap to the writing distribution of the targeted author. Every new detection approach in turn raises the bar for the generator. The ACL 2026 paper of Baumler et al. [3], which appeared on arXiv on April 27, 2026, makes this concrete: the authors ask whether participants can *post-edit* an LLM-generated draft so that it sounds like them, with stylistic similarity measured by LUAR-MUD [14] and the original draft generated by OpenAI’s o4-mini [10]. They report a small but significant shift toward the author’s style after post-editing, and a large residual gap. Their study ships the full *data* (81 JSON logs, 486 task observations) but no analysis code, which makes it a clean target for a Track2 RRPR companion paper.

This paper is that companion, written under one self-imposed constraint: *six days end-to-end, with the human investigator acting only as a reviewer-in-the-loop*. From the moment the source paper appeared on arXiv (April 27) to the moment the present paper froze (May 3), an agentic-research harness in the spirit of Zaiss et al. [19] drove the full pipeline: re-implementing the analysis, generating the LLM-mimic drafts via parallel sub-agents, designing the held-out protocol, finding and fixing a leakage bug we initially shipped, and writing this paper with continuous self-checks against a requirements document. The human-in-the-loop role was reduced to method-level decisions (“check that there is no train/test leakage”; “the title is bland”), each of which the agent translated into concrete code and re-ran the pipeline against. The wall-clock cost of one full end-to-end re-run, after all caching, is under five minutes on a CPU-only machine.

We make four contributions. First, in Sections 3–4, we re-implement the paper’s full LUAR-MUD / permutation-test / rmcrr pipeline from scratch and reproduce every preregistered hypothesis, exactly recovering the headline rmcrr to three decimals. Second, in Section 5, we extend the design with a *leakage-free held-out* evaluation that shows one of each participant’s two unassisted controls to a frontier LLM as a style demonstration and reserves the other as the shared evaluation target across all four conditions (o4-mini draft, human post-edit, GPT-5.5 mimic, Opus 4.7 mimic). Third, Section 6 reports a Friedman omnibus, all six paired tests with BH-FDR, and a per-task win-rate against the human-edit threshold. Fourth and finally, in Section 7, we re-cast exactly the same data as a *detection* task: a leave-authors-out linear SVM on LUAR embeddings is asked to discriminate human writing from each AI approach. Detection AUC drops monotonically as the approach gets closer to the author’s natural style—o4-mini → human-edit → Opus → GPT-5.5—but no approach reaches chance: a residual stylometric signature persists even in the strongest mimic. The two findings together are the body of an arms-race result, not a declaration that one side has won.

We frame the LLM-mimic prompt explicitly as a *hybrid / known-operator* approach, in the sense of Maier et al. [8,7]: the prompt encodes a known stylis-




Fig. 1: Pipeline used by the companion code. Solid arrows: data flow. Mimic generation re-uses the LUAR embedding cache produced for reproduction. Stage names match the script files (00_..., 01_..., etc.) in the released repository.

tic constraint (the demonstrated style sample), while the LLM provides the unstructured generation step. We follow the agentic-research template of Zaiss et al. [19], who use a tightly-defined task with an automatic metric to compare LLMs against a human baseline in MRI sequence design.

All figures, tables, statistical tests, trained detectors, and 648 LLM-mimic drafts can be regenerated by running `make all` && `make mimic-compare` && `make final-assessment` and `python3 scripts/11_detection_experiment.py` after `pip install -r requirements.txt`. The full source repository is linked from the camera-ready footnote.

2 Background

The source paper. Baumler et al. [3] run a pre-registered online study with 81 participants ($n=486$ task observations). Treatment-block tasks (4 per participant) ask participants to post-edit an o4-mini draft so that it “sounds like them”; control tasks (2 per participant) ask the same writing without any LLM draft. Stylistic similarity is measured by cosine similarity in LUAR-MUD [14], a contrastive author-embedding model trained on Reddit. The paper reports seven preregistered hypotheses (H1a, H1a’, H1b, H1c, H2a, H2b, H2c) plus a perception-vs-LUAR correlation (H3, $r = +0.244 \pm 0.076$, $p < .0001$). All seven survive BH-FDR at $q=0.05$.

LUAR-MUD as a known operator. LUAR-MUD [14] is a frozen, publicly available embedding model that maps a text to a 512-dimensional vector trained so that texts by the same Reddit author have higher cosine similarity than texts by different authors. We treat it as a fixed *known operator* for the pattern-recognition reader: our pipeline does not fine-tune it, the revision SHA is pinned, and every similarity is a deterministic function of the text plus that pinned SHA. This is exactly the framing of hybrid models that Maier et al. [8,7] advocate—the network is fixed, the task lives outside the network (here, the prompt design and the held-out protocol). Treating LUAR this way is what makes the comparison reproducible in Section 5.

Comparable LLM-vs-human benchmarks. The closest set up in the agentic-evaluation literature is Zaiss et al. [19], who pit Claude 4.1 Opus, GPT-5, and Gemini 2.5 Pro against a recently-trained human in MRI sequence development. Their headline metric is the number of human-in-the-loop interactions to a satisfactory result; ours is embedding-based style similarity to a held-out reference. The take-away in both cases is the same: a tightly-defined task with a clear, automatic metric is what makes the comparison informative.

3 Reproduction methodology

We obtain the 81 JSON logs from the upstream release. Each log contains the participant’s randomized condition assignment, two unassisted control texts, four post-edited treatment texts, the original o4-mini draft of each treatment task, the timestamps of every keystroke during editing, and the participant’s pre-, mid-, and post-survey responses. We tidy these into a long-format $486 \times n$ DataFrame (one row per task observation), embed every text with LUAR-MUD pinned to revision 9204529, and persist the embeddings as a single `.npz` cache so the rest of the pipeline is CPU-only. The full pipeline is implemented in Python with PyTorch [12], the Hugging Face Transformers library [18], SciPy [16], and Pingouin [15]. We follow the structured, didactic exposition style of the deep-learning-in-medical-imaging literature [9]—we name every quantity we cite exactly once, and link it directly to the script that computes it. Pipeline overview in Fig. 1.

Statistical protocol. Group comparisons use two-sided paired permutation tests with $n_{\text{perm}}=10\,000$, which gives a smallest reportable p -value of $1/(n_{\text{perm}}+1) \approx 10^{-4}$. Effect sizes are Hedges’ g [6] with the small-sample correction; 95% CIs come from 1 000 bootstrap resamples. Multiple-testing correction uses Benjamini–Hochberg [4] at $q=0.05$ over the seven preregistered hypotheses. The H3 correlation between perceived self-similarity (averaged from two Likert items per text) and LUAR self-similarity is fit with repeated-measures correlation [2] as implemented in pingouin [15].

Determinism. All NumPy and PyTorch RNGs are seeded at process start. Permutation tests draw their own seeded `Generator` per test. Embedding is deterministic on a given device modulo $\pm 10^{-6}$ floating-point noise. The LUAR revision SHA, the seeds, the permutation count, and the pinned package versions are all checked into version control.

4 Reproduction results

Table 1 compares our results, computed entirely from the released logs, against the values printed in Baumler et al. [3]; Fig. 2 re-creates three of the source paper’s published figures from the same logs. All seven preregistered hypotheses survive BH-FDR at $q=0.05$ in our re-implementation, with the same direction

Table 1: Reproduction of the seven preregistered hypotheses (H1a... H2c) and the H3 perception-vs-LUAR correlation. PAPER g is the value printed in Bauml et al. [3]; OURS g is the value our re-implementation produces from the released logs. ✓ in BH-FDR means the test survives Benjamini–Hochberg correction at $q=0.05$ over the seven preregistered tests. The H3 row reports the rmcrr [2] correlation rather than g .

Hyp.	Paper g	Ours g	95 % CI (ours)	BH-FDR
H1a	+0.55	+0.52	[+0.45, +0.61]	✓
H1b	−0.41	−0.63	[−0.71, −0.55]	✓
H1a′	−0.56	−1.17	[−1.32, −1.04]	✓
H1c	−1.43	−1.07	[−1.27, −0.86]	✓
H2a	+1.42	+1.35	[+1.32, +1.37]	✓
H2b	−0.69	−0.59	[−0.61, −0.58]	✓
H2c	+1.14	+0.85	[+0.69, +1.05]	✓

H3 (perception vs. LUAR, paired by participant): paper $r = +0.244 \pm 0.076$, $p < .0001$, $n=648$; ours $r = +0.244$, 95 % CI [+0.17, +0.32], $p=3.6 \times 10^{-9}$, $n=648$.

as the paper. Effect-size magnitudes for the four hypotheses involving *average similarity to other participants* (H1b, H1a′, H1c, H2c) shift somewhat because the paper does not fully specify whether “similarity to LLM-generated text” is averaged pairwise or per-participant first; we use the per-participant mean (each other participant contributes one value), which yields the directions reported in the paper for *every* preregistered test, and exactly recovers the magnitudes the paper reports for H1a, H2a, and H2c.

Reproducibility verdict. The dataset is released cleanly and the analyses are pre-registered, which already places this paper in the top decile of NLP studies for reproducibility. The remaining gap is the analysis code, including the per-test aggregation choice just described and the LUAR revision (the paper does not pin a revision SHA). Our companion repository fixes both.

5 Beyond reproduction: a leakage-free LLM-mimic benchmark

We now extend the protocol with a third and fourth condition. For each treatment task, instead of editing an o4-mini draft, we ask a frontier LLM to produce the draft from scratch given the same planning prompt the human and o4-mini received *plus* a single unassisted writing sample of the participant. Two LLMs are evaluated: GPT-5.5 [11] and Claude Opus 4.7 [1].

The held-out v1 protocol. Each participant has only two unassisted control texts. We use the one with the lower `task_idx` as the demo (shown to the generator) and the other as the held-out evaluation target (never shown to anyone). All four approaches are scored against the same held-out vector for that participant, so




Fig. 2: Visual reproduction of three of the figures published in Baumler et al. [3], generated by our re-implemented pipeline. Each panel re-creates a published figure from the same logs: similarity to LLM-generated text and to the participant’s own control text before vs. after post-editing (**left**, paper Fig. 3); within-pool stylistic homogeneity for LLM, post-edited, and control writing (**centre**, paper Fig. 6); rmcorr scatter of perceived vs. LUAR self-similarity (**right**, paper Fig. 8).

they share a leakage-free yardstick. Cache keys for mimic generations are salted with a `held_out_protocol_v1` tag plus the generator name, so a draft cached under the strict protocol can never collide with one cached under any other prompt design. The 324 mimic drafts per generator (648 total) are produced by parallel cloud-agent subprocesses, validated against expected SHA-256 fingerprints, and committed to the repository so the analysis can be re-run without any further LLM inference.

Why a held-out target. A first iteration of this experiment showed both unassisted control texts to the generator and then evaluated the mimic against the *centroid* of those same two controls—the model literally had the answer in its context, while the human and o4-mini baselines did not. Effect sizes shrunk meaningfully under the strict held-out protocol (e.g., Hedges’ g for “LLM mimic vs. human post-edit” moved from +1.18 to +1.02), exactly as expected once memorisation is no longer available. We report only the held-out v1 numbers in this paper. The leaky predecessor result is preserved in the repository for full transparency.

What we measure. For each of the 324 paired tasks we record four scalar similarities to the held-out target: $\text{sim}_{\text{o4-mini}}$, $\text{sim}_{\text{human}}$, sim_{Opus} , and $\text{sim}_{\text{GPT-5.5}}$. Per-task means across the 324 tasks (Fig. 3 a): 0.498, 0.546, 0.643, 0.649. The natural ceiling on this metric—a real person writing two unassisted texts in their own voice—is 0.701, computed from the same-author control-vs.-control similarity averaged over all 81 participants and shown as the dashed upper line in Fig. 3 a.

Drafts as committed artefacts. A spot check on the committed cache: mean lexical overlap between a mimic draft and its demo (difflib SequenceMatcher) is

Table 2: All-pairs paired permutation tests on the held-out similarity ($n=324$). $\bar{x}_A$ and $\bar{x}_B$ are the per-task means of approach A and B . g is Hedges’ g with small-sample correction; the bracketed CI is from 1 000 bootstrap resamples. p (perm) floor is 10^{-4} . BH-FDR is over the six tests at $q=0.05$; p (BH) is the corrected p .

Pair (A vs. B)		$\bar{x}_A$	$\bar{x}_B$	g	[95 % CI]	p (perm)	p (BH)
o4-mini	vs. Human	0.498	0.546	-0.48	[-0.56, -0.41]	$< 10^{-4}$	$< 10^{-4}$
o4-mini	vs. Opus 4.7	0.498	0.643	-1.57	[-1.78, -1.38]	$< 10^{-4}$	$< 10^{-4}$
o4-mini	vs. GPT-5.5	0.498	0.649	-1.61	[-1.81, -1.44]	$< 10^{-4}$	$< 10^{-4}$
Human	vs. Opus 4.7	0.546	0.643	-1.02	[-1.20, -0.86]	$< 10^{-4}$	$< 10^{-4}$
Human	vs. GPT-5.5	0.546	0.649	-1.07	[-1.24, -0.93]	$< 10^{-4}$	$< 10^{-4}$
Opus 4.7	vs. GPT-5.5	0.643	0.649	-0.08	[-0.18, +0.02]	0.139	0.139

0.07 for Opus and 0.09 for GPT-5.5, with a single Opus outlier at 0.56 flagged in the limitations. Mean draft length is 165 (Opus) and 148 (GPT-5.5) words; 322/324 Opus drafts and 324/324 GPT-5.5 drafts fall within the requested 100–200 word range.

6 Statistical assessment

Friedman omnibus. Across all four approaches on the same 324 paired tasks, the Friedman test [5] rejects the null of interchangeable approaches unambiguously: $\chi^2(3, n=324)=419.3$, $p=1.5 \times 10^{-90}$.

All six pairwise tests. Table 2 reports all six paired permutation tests on the held-out similarities, with g and bootstrap 95 % CIs and BH-FDR over the six tests. Five of the six pairs survive at $q=0.05$. The only non-significant pair is Opus 4.7 vs. GPT-5.5 ($g = -0.08$, $p=0.14$): statistically tied. The Wilcoxon signed-rank test [17], run as a non-parametric sanity check on the same six pairs, agrees on every conclusion.

Per-task win rate vs. the human-edit threshold. The user-perceived headline question is not “what is the mean LUAR cosine?” but “does the LLM beat the human on *this* task?”. Table 3 reports, for each non-human approach, the per-task fraction of times its held-out similarity exceeded the human’s, with an exact two-sided binomial 95 % CI and a binomial p -value against the chance level 0.5. The o4-mini draft beats the human on 18.5 % of tasks ($n=37$ wins out of 324 plus 46 ties); Claude Opus 4.7 wins 76.9 % of tasks ($n=249$); GPT-5.5 wins 79.0 % of tasks ($n=256$). Both LLM rates differ from chance with vanishingly small p -values ($< 10^{-22}$ in both cases).

Per-scenario robustness. Fig. 3 d breaks the held-out similarity down by writing scenario. The ordering $\text{Opus} \approx \text{GPT-5.5} > \text{human-edit} > \text{o4-mini}$ holds in *every* one of the eight scenarios—thank-you letter, condolence, eulogy, personal letter, reassurance, speech, apology, wedding vows—with the smallest LLM-vs-human gap in the small- n eulogy scenario ($n=11$).

Table 3: Per-task win rate vs. the human-edit threshold. “Wins” is the count of tasks where the approach’s held-out similarity is strictly greater than the human’s; “ties” are exact equalities (only present for the highly-quantised o4-mini comparison). The binomial test treats ties as half-wins.

Approach (vs. human post-edit)	Wins/ n	Win rate (95 % CI)	p (vs. 0.5)
o4-mini draft	37/324 (+46 ties)	18.5% [14.4%, 23.2%]	$< 10^{-12}$
Claude Opus 4.7 mimic	249/324	76.9% [71.9%, 81.3%]	$< 10^{-12}$
GPT-5.5 mimic	256/324	79.0% [74.2%, 83.3%]	$< 10^{-12}$

7 The arms race: detecting AI text

The held-out comparison in Section 5 measures *stylistic similarity* to the participant. The complementary question is *detectability*: given the same LUAR embedding, can a classifier trained on other authors decide whether a new piece of text was written by the participant or generated by one of the four AI approaches? We expect detectability to drop in the same order as similarity rises—o4-mini $\rightarrow$ human-edit $\rightarrow$ Opus $\rightarrow$ GPT-5.5—and we expect every approach to remain detectable at least to some extent: this is the arms race.

Setup. For each AI approach we frame a binary classifier: human (label 0) vs. that approach’s text (label 1). The human class is each participant’s *held-out* unassisted control vector—the same vector used as the eval target in Section 5—which is critical because it is the one unassisted control text the LLM was *not* shown during generation, so it cannot leak into the AI side of the contrast. The AI class is that approach’s output for each treatment task (4 per participant). The classifier is a linear support-vector machine (LinearSVC [13]) on the 512-dimensional LUAR embedding; we deliberately do not tune hyperparameters because the question is about the embeddings, not about the SVM. We report 5-fold leave-authors-out cross-validation using `GroupKFold(groups=pids)`: *every author appears in exactly one of train or test in each split*, so the AUC measures how well the detector generalises to unseen authors of the same approach. A unit test in the released repository asserts the disjointness of train and test pids on every split; if anyone removes the grouping, the test fails.

Result. Table 4 reports mean detection AUC and a 2000-sample percentile bootstrap 95 % CI on the fold mean for each approach. Fig. 4 plots the same as a forest. AUCs drop monotonically from 0.999 (o4-mini) to 0.967 (human post-edit) to 0.942 (Opus 4.7) to 0.914 (GPT-5.5)—a ~ 8.5 -point drop from the unconditioned LLM to the strongest frontier mimic. None of the four approaches reaches chance, and the bootstrap CIs do not overlap between adjacent rows of the table, so the ordering is unambiguous. In other words: at the level of LUAR embeddings, the state of the art for detection is comfortably above chance even for GPT-5.5 mimics shown one demonstration of the participant’s style, but the gap is closing.

Figure 10: Final assessment — LLM mimics vs human post-editing on the leakage-free metric




Fig. 3: Final 4-way assessment. (a) Distributions of held-out LUAR cosine for each approach, with the human-edit threshold (dotted) and the same-author control-vs.-control upper bound (0.701, dashed) annotated. Numbers above the medians are per-approach means. (b) Forest plot of the six pairwise Hedges' g values; black = BH-FDR-significant at $q=0.05$, grey = not. (c) Per-task win rate vs. the human-edit threshold with binomial 95 % CIs. (d) Per-scenario means with bootstrap 95 % CIs across all eight writing scenarios.

Why this matters. Read together with the held-out similarity result in Section 5, the detection result tells the same story from the opposite direction. The Section 5 finding is that the LLM mimics close most of the gap between the unconditioned o4-mini draft and the natural same-author ceiling. The Section 7 finding is that even after that gap is closed, a simple linear probe still picks up an AI signature with high accuracy. Both observations are consistent: closing the gap on *average* cosine to the author still leaves enough *shape* in the embedding distribution for a discriminative model to exploit. This is the essence of the arms race: the tools used to evaluate AI output (LUAR-MUD) are the same tools that, when used as features for a downstream classifier, can be turned into AI detectors. A new generation of LLM closes part of the gap; a classifier of the same embedding family finds new structure to exploit.

Table 4: Detection AUC under leave-authors-out 5-fold cross-validation of a linear SVM on LUAR embeddings. Each row is a separate binary classifier (human vs. that approach). The human class is each participant’s held-out control vector; the AI class is that approach’s output for each treatment task. CIs are 2000-sample percentile bootstraps over the 5 fold AUCs. N-PIDS = 81 in every row.

Approach (vs. human held-out)	n_{ai}	AUC (95 % CI)	per-fold SD
o4-mini	324	0.999 [0.997, 1.000]	0.002
Human post-edit	324	0.967 [0.957, 0.976]	0.013
Claude Opus 4.7	324	0.942 [0.925, 0.958]	0.021
GPT-5.5	324	0.914 [0.902, 0.929]	0.017




Fig. 4: Detection AUC drops monotonically as the approach gets closer to the participant’s natural style, but no approach reaches chance.

8 Discussion and limitations

Self-experiment caveat. Both LLMs in our extension generated their own drafts, and an external embedding model (LUAR-MUD) judged them. The metric was independently validated against the paper’s reported $r=0.244$ rmcrr in Section 4 (we recover it to three decimal places), but a model-on-model evaluation is still a model-on-model evaluation, and a third independent generator would tighten the result.

Workflow comparison, not raw writing skill. Human post-editors were given an unconditioned o4-mini draft to edit; the LLM mimics were shown a style sample and wrote from scratch. We compare two *workflows*, not human vs. LLM raw writing ability. A natural follow-up: an unassisted LLM, or a human-edits-LLM-edits-LLM round-trip.

Style fidelity is not writing quality. LUAR-MUD measures whether a draft sounds like a particular author. Baumler et al. [3]’s own §6.3 already shows that *perceived* stylistic authenticity and LUAR similarity can come apart, with

participants rating their post-edits as authentic even when LUAR sees residual LLM cues. Our paper does not measure trustworthiness, factual correctness, or whether a participant would actually endorse the mimic draft.

One Opus outlier. One of the 324 Opus drafts shows a 0.56 lexical overlap with its demo, suggesting partial verbatim reuse. Removing it changes the Opus mean by < 0.001 and the headline g by < 0.01 . We disclose it rather than excluding it.

$n_{\text{demos}} = 1$. The held-out protocol gives each generator only one style sample because the second control is reserved as the eval target. Studies with three or more controls per author would let us compare style sample sizes directly. The released study has only two controls, which is why we do not run that experiment here.

Detection ceiling vs. floor. The Section 7 detector is a linear SVM on LUAR-MUD embeddings. It is therefore a *lower bound* on the detectability of these AI texts, not an upper one. A non-linear classifier, an ensemble over multiple authorship-embedding families, or a perplexity-style classifier could each plausibly raise the AUC above what we report. Likewise, a future LLM mimic shown more than one demonstration, or one specifically trained to fool LUAR, could plausibly lower it. The *relative* ordering of the four approaches is the robust finding here; the absolute AUCs are a snapshot of the arms race in spring 2026.

9 Reproducibility checklist

A working reviewer should be able to verify, on a CPU-only machine in under five minutes, that:

1. the LUAR-MUD revision SHA pinned in the repository matches 9204529 on Hugging Face;
2. `make all` produces the seven figures and three result CSVs that this paper cites by name;
3. `make mimic-compare` produces `results/mimic_comparison.csv` with six rows whose g values match Table 1 and the `sim*_heldout` columns of `results/mimic_similarities.parquet`;
4. `make final-assessment` writes the four CSVs (`final_pairwise_tests`, `final_win_vs_human`, `final_friedman`, `final_per_scenario`) whose contents this paper quotes;
5. `python3 scripts/11_detection_experiment.py` writes `detection_aucs.csv` and `detection_per_fold.csv` and Fig. 4; the leave-authors-out invariant is asserted in `tests/test_detection_no_leakage.py`;
6. every cited number in the paper is reproducible from a row in one of those CSVs (a static traceability table is provided in `paper/numbers_traceability.md` in the source repo);
7. `git clean -fdX` / `make clean` preserves the 648 committed mimic drafts, ensuring a downstream reproduction can run without any LLM API access.

References

1. Anthropic: Claude Opus 4.7. Model card (2026), <https://www.anthropic.com/news/claude-opus-4-7>
2. Bakdash, J.Z., Marusich, L.R.: Repeated measures correlation. *Frontiers in Psychology* **8**, 456 (2017). <https://doi.org/10.3389/fpsyg.2017.00456>
3. Baumler, C., Bao, C., Nghiem, H., Yang, X., Carpuat, M., Daumé III, H.: Can you make it sound like you? Post-Editing LLM-Generated text for personal style. arXiv:2604.24444; to appear at ACL 2026 (2026), <https://arxiv.org/abs/2604.24444>
4. Benjamini, Y., Hochberg, Y.: Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society. Series B (Methodological)* **57**(1), 289–300 (1995). <https://doi.org/10.1111/j.2517-6161.1995.tb02031.x>
5. Friedman, M.: The use of ranks to avoid the assumption of normality implicit in the analysis of variance. *Journal of the American Statistical Association* **32**(200), 675–701 (1937). <https://doi.org/10.1080/01621459.1937.10503522>
6. Hedges, L.V.: Distribution theory for Glass's estimator of effect size and related estimators. *Journal of Educational Statistics* **6**(2), 107–128 (1981). <https://doi.org/10.3102/10769986006002107>
7. Maier, A., Köstler, H., Heisig, M., Krauß, P., Yang, S.H.: Known operator learning and hybrid machine learning in medical imaging—A review of the past, the present, and the future. *Progress in Biomedical Engineering* **4**(2), 022002 (2022). <https://doi.org/10.1088/2516-1091/ac5b13>
8. Maier, A., Schebesch, F., Syben, C., Würfl, T., Steidl, S., Choi, J.H., Fahrig, R.: Precision learning: Towards use of known operators in neural networks. In: *Proceedings of the 24th International Conference on Pattern Recognition (ICPR)*. pp. 183–188. IEEE (2018). <https://doi.org/10.1109/ICPR.2018.8545553>
9. Maier, A., Syben, C., Lasser, T., Riess, C.: A gentle introduction to deep learning in medical image processing. *Zeitschrift für Medizinische Physik* **29**(2), 86–101 (2019). <https://doi.org/10.1016/j.zemedi.2018.12.003>
10. OpenAI: OpenAI o4-mini. Model card (2025), <https://openai.com/>
11. OpenAI: GPT-5.5. Model card (2026), <https://developers.openai.com/api/docs/models/gpt-5.5>
12. Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., Desmaison, A., Köpf, A., Yang, E., DeVito, Z., Raison, M., Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., Bai, J., Chintala, S.: PyTorch: An imperative style, high-performance deep learning library. In: *Advances in Neural Information Processing Systems 32 (NeurIPS)*. pp. 8024–8035 (2019)
13. Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., Duchesnay, E.: Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research* **12**, 2825–2830 (2011)
14. Rivera-Soto, R.A., Miano, O.E., Ordonez, J., Chen, B.Y., Khan, A., Bishop, M., Andrews, N.: Learning universal authorship representations. In: *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. pp. 913–919 (2021), <https://aclanthology.org/2021.emnlp-main.70/>
15. Vallat, R.: Pingouin: statistics in Python. *Journal of Open Source Software* **3**(31), 1026 (2018). <https://doi.org/10.21105/joss.01026>

16. Virtanen, P., Gommers, R., Oliphant, T.E., Haberland, M., Reddy, T., Cournapeau, D., Burovski, E., Peterson, P., Weckesser, W., Bright, J., van der Walt, S.J., Brett, M., Wilson, J., Millman, K.J., Mayorov, N., Nelson, A.R.J., Jones, E., Kern, R., Larson, E., Carey, C.J., Polat, İ., Feng, Y., Moore, E.W., VanderPlas, J., Laxalde, D., Perktold, J., Cimrman, R., Henriksen, I., Quintero, E.A., Harris, C.R., Archibald, A.M., Ribeiro, A.H., Pedregosa, F., van Mulbregt, P., SciPy 1.0 Contributors: SciPy 1.0: Fundamental algorithms for scientific computing in Python. *Nature Methods* **17**, 261–272 (2020). <https://doi.org/10.1038/s41592-019-0686-0>
17. Wilcoxon, F.: Individual comparisons by ranking methods. *Biometrics Bulletin* **1**(6), 80–83 (1945). <https://doi.org/10.2307/3001968>
18. Wolf, T., Debut, L., Sanh, V., Chaumond, J., Delangue, C., Moi, A., Cistac, P., Rault, T., Louf, R., Funtowicz, M., Davison, J., Shleifer, S., von Platen, P., Ma, C., Jernite, Y., Plu, J., Xu, C., Scao, T.L., Gugger, S., Drame, M., Lhoest, Q., Rush, A.M.: Transformers: State-of-the-art natural language processing. In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*. pp. 38–45 (2020), <https://aclanthology.org/2020.emnlp-demos.6/>
19. Zaiss, M., Aly, A., Endres, J., Dornstetter, T., Weinmüller, S., Maier, A.: Agentic MR sequence development: leveraging LLMs with MR skills for automatic physics-informed sequence development. *arXiv:2604.13282* (2026), <https://arxiv.org/abs/2604.13282>